

CORD-19 Explore dataset

In general, this jupyter notebook is designated to explore the CORD19 dataset:

Wade, Alex D.; Williams, Ivana (2021), CORD-19 Software Mentions, Dryad, Dataset, <https://doi.org/10.5061/dryad.vmcvdcns0> (<https://doi.org/10.5061/dryad.vmcvdcns0>).

First, relevant packages must be imported into the notebook.

In [1]:

```
import numpy as np
import pandas as pd
import csv
import ast
import collections
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import Levenshtein as lev
from fuzzywuzzy import fuzz
import datetime
import matplotlib.pyplot as plt
import re
from urllib.parse import urlparse
from collections import Counter
```

Get the data and save it to a variable.

In [2]:

```
CORD19_CSV = pd.read_csv('../data/cord-19/CORD19_software_mentions.csv')
```

Show the head of the dataset to inspect all columns and obtain a broad overview.

In [3]:

```
CORD19_CSV
```

Out[3]:

	paper_id	doi	title	source_x	license
0	00006903b396d50cc0037fed39916d57d50ee801	NaN	Urban green space and happiness in developed c...	ArXiv	a
1	0000fcce604204b1b9d876dc073eb529eb5ce305	10.1016/j.regg.2021.01.002	La Geriatria de Enlace con residencias en la é...	Elsevier; PMC	cc
2	000122a9a774ec76fa35ec0c0f6734e7e8d0c541	10.1016/j.rec.2020.08.002	Impact of COVID-19 on ST-segment elevation myo...	Elsevier; Medline; PMC	nc
			Absence of surface	Elsevier;	

General check of the DataFrame

First of all, it is interesting to investigate how many NaNs are contained within the DataFrame.

In [4]:

```
CORD19_CSV.isnull().sum().sum()
```

Out[4]:

12476

How are the NaNs distributed between the columns of the DataFrame?

In [5]:

```
def count_nan_per_column(df):  
    """  
    This function counts the NaNs per column of the received DataFrame  
    and returns a DataFrame containing a column listing the number of NaNs..  
    """  
    col_num = []  
    col_name = []  
    for i in df:  
        col_num.append(df[i].isnull().sum())  
    col_names = df.columns  
    df_joined = pd.DataFrame(col_num, col_names)  
    df_joined.columns = ['NaNs']  
    return df_joined
```

In [6]:

```
df_NaNs = count_nan_per_column(CORD19_CSV)
df_NaNs
```

Out[6]:

	NaNs
paper_id	0
doi	3144
title	0
source_x	0
license	0
publish_time	0
journal	9332
url	0
software	0

For further analysis, it is important to obey that the columns "doi" and "journal" contain NaNs.

In [7]:

```
ser_col_length = CORD19_CSV.fillna('').astype(str).apply(lambda x:x.str.len()).mean()
ser_col_length_sorted = ser_col_length.sort_values(ascending=False)
ser_col_length_sorted
```

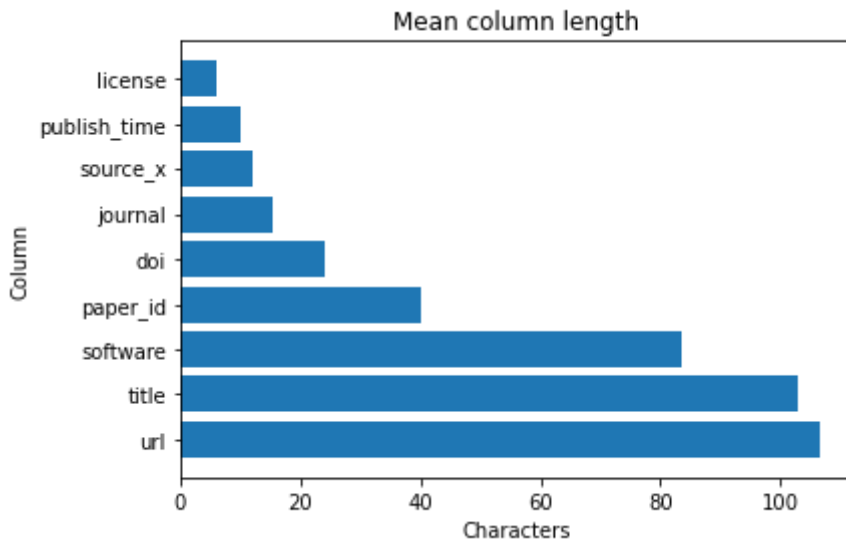
Out[7]:

```
url          106.520155
title        102.960064
software      83.598324
paper_id     40.000000
doi          24.133883
journal       15.316212
source_x     12.139758
publish_time  9.930508
license       5.899985
dtype: float64
```

In [8]:

```
x = ser_col_length_sorted.index
y = ser_col_length_sorted.values

plt.barh(x, y)
plt.title("Mean column length")
plt.xlabel("Characters")
plt.ylabel("Column")
plt.show()
```



For further analysis, the average length of titles and URLs is long and can be seen as descriptive.

The column `publish_time` is converted to the `DateTime` format.

In [9]:

```
CORD19_CSV['publish_time'] = pd.to_datetime(CORD19_CSV.publish_time)
```

Dedicated analysis based on columns

Create own series for each column.

In [10]:

```
paper_id = CORD19_CSV.paper_id
doi = CORD19_CSV.doi
title = CORD19_CSV.title
source_x = CORD19_CSV.source_x
license = CORD19_CSV.license
publish_time = CORD19_CSV.publish_time
journal = CORD19_CSV.journal
url = CORD19_CSV.url
```

Checking the column "paper_id"

First, check the integrity of the series `paper_id`. It should be obvious that each row has a unique ID that is not shared with another row.

In [11]:

```
paper_id_counted = paper_id.value_counts()
paper_id_counted
```

Out[11]:

```
0ed3c6a5559cd73307184f51fc53ccc76da559bc    3
5d6678f81812464543b367e7de138e23b3483ed1    2
5d0d0bd116976e1412c10a84902894999df4a342    2
ff40e6b44e151e42a54227e255a88d0c0c104876    2
46b053c7126c1603101f46e4bb6e411f790a45fc    2
..
8193755d869fc3ecf58f4838fcd616163584cbdc    1
74165e4814a803cf72f62b9cad68c00659607c1a    1
ba32b2acc5045255ab8b9d17083b1196131dd700    1
9c1595012ce2d5e518584750984dfbb69771e0b5    1
f25f980338756b6535ec83c06ae483d56af3e085    1
Name: paper_id, Length: 77436, dtype: int64
```

The first finding linked to the column `paper_id` is that this column does not solely contain unique ID's. Some rows share ID's. For this, the exact number of shared ID's need to be found.

In [12]:

```
def find_shared_values(col):
    """ This function checks columns for shared values and deducts NaNs
    """
    col_shared = len(col) - len(col.value_counts())
    col_shared = col_shared - col.isnull().sum()
    return col_shared
```

In [13]:

```
paper_id_shared_ids_num = find_shared_values(paper_id)
paper_id_shared_ids_num
```

Out[13]:

The column contains 12 shared ID's with even one used three times.

In [14]:

```
paper_id_counted.head(paper_id_shared_ids_num)
```

Out[14]:

```
0ed3c6a5559cd73307184f51fc53ccc76da559bc    3
5d6678f81812464543b367e7de138e23b3483ed1    2
5d0d0bd116976e1412c10a84902894999df4a342    2
ff40e6b44e151e42a54227e255a88d0c0c104876    2
46b053c7126c1603101f46e4bb6e411f790a45fc    2
dd74a3a343529174fe7c6485723cf2d5911c18ed    2
d1dde1df11f93e8eae0d0b467cd0455afdc5b98c    2
0831fe32280e46ba8d5c1a9456111e1e009863ac    2
ec7d3038b8912a9fc92f4d02a2c30d566d4d0a93    2
36e2047d1674c3095617f3eb97f9f61e48989dfe    2
c89f86cdd9d41eeec127cc0b03990c52888a9635    2
b391f95092b4335dc9f80fa3a1d56ff9e1d3f8bf    1
Name: paper_id, dtype: int64
```

Check if the dataset contains only duplicates for the column `paper_id` or whole DataFrame rows. Therefore, the function `collect_rows_of_df` will support the process.

In [15]:

```
#Method receiving string and dataframe which returns double or tripple dataframe to append
def collect_rows_of_df(df,column,st):
    """This function receives a dataframe, a column contained within a dataframe
    and a string which can be found within the column.
    Then, the string is compared to the whole column and.
    When a match is found, the corresponding rows are returned as a dataframe.
    """
    subset = df[df[column] == st]
    return subset
```

Collecting rows that share their `paper_id`.

In [16]:

```
x = 0
shared_paper_id_df = pd.DataFrame(columns=['paper_id', 'doi', 'title', 'source_x', 'license', 'p
while x < paper_id_shared_ids_num:
    shared_paper_id_df = shared_paper_id_df.append(collect_rows_of_df(CORD19_CSV, 'paper_id
    x= x+1
shared_paper_id_df
```

Out[16]:

	paper_id	doi	title	source_x	license
4466	0ed3c6a5559cd73307184f51fc53ccc76da559bc	10.1016/j.jinf.2020.02.019	Simulating and forecasting the cumulative conf...	Elsevier; Medline; PMC	no-cc
4467	0ed3c6a5559cd73307184f51fc53ccc76da559bc	10.1016/j.jinf.2020.02.020	Novel coronavirus disease (Covid-19): The firs...	Elsevier; Medline; PMC	els-covid
4468	0ed3c6a5559cd73307184f51fc53ccc76da559bc	10.1016/j.jinf.2020.02.011	Chinese medical personnel	Elsevier; Medline;	els-covid

Besides software, the paper_id duplicates have variations among all other columns which means that the same paper_id was used for different publications.

Checking the column "doi"

Each paper in this dataset should have a unique "doi" because duplicates are skewing the analysis and potential outcomes.

In [17]:

```
doi_counted = doi.value_counts()
doi_counted
```

Out[17]:

```
10.1016/j.dsx.2020.04.012      2
10.31729/jnma.5498             2
10.1097/im9.0000000000000035    1
10.1016/j.jmoldx.2014.12.002    1
10.1007/978-3-030-52237-7_31    1
..
10.1016/j.jairtraman.2020.101900 1
10.1101/2020.07.16.20155721      1
10.1093/cid/ciaa468             1
10.1016/j.jtumed.2020.06.005    1
10.1007/s12117-020-09401-y      1
Name: doi, Length: 74302, dtype: int64
```

In [18]:

```
doi_shared_dois_num = find_shared_values(doi)
doi_shared_dois_num
```

Out[18]:

2

The column "doi" contains two entries that share a "doi".

In [19]:

```
doi_counted.head(doi_shared_dois_num)
```

Out[19]:

```
10.1016/j.dsx.2020.04.012      2
10.31729/jnma.5498             2
Name: doi, dtype: int64
```

Based on rows, it needs to be explored how the shared doi's are affecting each other.

In [20]:

```
x = 0
shared_doi_df = pd.DataFrame(columns=['paper_id', 'doi', 'title', 'source_x', 'license', 'publis
while x < doi_shared_dois_num:
    shared_doi_df = shared_doi_df.append(collect_rows_of_df(CORD19_CSV, 'doi', doi_counted
    x = x+1
shared_doi_df
```

Out[20]:

	paper_id	doi	title	source
16564	36e2047d1674c3095617f3eb97f9f61e48989dfe	10.1016/j.dsx.2020.04.012	Artificial Intelligence (AI) applications for ...	PI
16565	36e2047d1674c3095617f3eb97f9f61e48989dfe	10.1016/j.dsx.2020.04.012	Artificial Intelligence (AI) applications for ...	Elsev
38600	80273c63683cad57323802542cfdcfcd76c805bf	10.31729/jnma.5498	Mental Wellbeing during the Lockdown Period fo...	PI
42407	8cab7532249cedf3815d4dada6400390a1f8a28a	10.31729/jnma.5498	Interpersonal Violence during the COVID-19 Loc...	PI

The first matched doi is already known for sharing paper_id's. For the second shared doi, there is an interesting

occurrence because the algorithm which created the dataset created two distinct outcomes for the column "software". The affected entries are presented below.

In [21]:

```
corrupted_software_mentions = pd.DataFrame()  
corrupted_software_mentions['Publication 1'] = shared_doi_df.loc[38600]  
corrupted_software_mentions['Publication 2'] = shared_doi_df.loc[42407]  
corrupted_software_mentions
```

Out[21]:

	Publication 1		Public
paper_id	80273c63683cad57323802542cfdcfd76c805bf	8cab7532249cedf3815d4dada6400390a1	
doi	10.31729/jnma.5498	10.31729/jnr	
title	Mental Wellbeing during the Lockdown Period fo...	Interpersonal Violence during the CC	
source_x	PMC		
license	cc-by		
publish_time	2020-10-31 00:00:00	2020-10-31 0	
journal	JNMA J Nepal Med Assoc	JNMA J Nepal Mec	
url	https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7...	https://www.ncbi.nlm.nih.gov/pmc/articles/F	
software	['SPSS']		['SPSS',

In comparison to publication 1, publication 2 is affiliated additionally to the software mention "WHO". This finding leads to the assumption that the algorithm which is creating software mentions does not work in a reproducible way.

Checking the column "title"

Next, the column title will be investigated to produce insights.

In [22]:

```
title_counted = title.value_counts()
title_counted
```

Out[22]:

```
Full Issue PDF
15
Posters
8
Poster Presentations
8
Poster Sessions
7
Poster Session Abstracts
7

..
Population Data-Driven Formulation of a COVID-19 Therapeutic
1
Metformin Reduces NGF-Induced Tumour Promoter Effects in Epithelial Ovarian
Cancer Cells
1
Viral infection detection using metagenomics technology in six poultry farms
of eastern China
1
Effects of work status changes and perceived stress on glycaemic control in
individuals with type 1 diabetes during COVID-19 lockdown in Italy
1
STRATEGY FOR THE CONTAINMENT, MITIGATION, AND SUPPRESSION OF THE COVID-19 PA
NDEMIC IN FRAGILIZED COMMUNITIES ON THE PERIPHERY OF A LARGE BRAZILIAN CITY.
1
Name: title, Length: 76578, dtype: int64
```

In [23]:

```
title_shared_titles_num = find_shared_values(title)
title_shared_titles_num
```

Out[23]:

```
870
```

Compared to the other duplicates, there are many more shared titles.

In [24]:

```
if title_shared_titles_num > 20:  
    head_for_title = 40  
else:  
    head_for_title = title_shared_titles_num  
title_counted.head(head_for_title)
```

Out[24]:

```
Full Issue PDF  
15  
Posters  
8  
Poster Presentations  
8  
Poster Sessions  
7  
Poster Session Abstracts  
7  
Scientific Abstracts  
6  
Physicians Poster Sessions  
6  
Mitteilungen der Deutschen Gesellschaft für Neurologie  
5  
NEWS  
5  
Infectieziekten  
5  
Abstracts cont.  
5  
Abstract  
5  
Virtual Meeting Announcment  
4  
Mitteilungen der ÖGKJ  
4  
Tracing the Origins of Agricultural Products with Barcoded Microbial Spores  
4  
Oral Abstracts  
3  
Mitteilungen des BDI  
3  
Mitteilungen der DGN  
3  
China  
3  
The EU versus the Automotive Industry  
3  
Oral presentations  
3  
Poster presentations  
3  
Oral Presentations  
3  
Introduction to Computer Network Vulnerabilities  
3  
Report  
3  
ISAR News
```

```

3
Highlights From the Current Issue - Audiovisual Summary
3
Forum
3
Using Morpholinos to Control Gene Expression
3
Computer Network Vulnerabilities
3
Reply
3
New approximations, and policy implications, from a delayed dynamic model of
a fast pandemic      3
Prediction of repurposed drugs for treating lung injury in COVID-19
3
Mitteilungen der DGIM
3
Symposium Summaries
3
Bildgebende Verfahren: Röntgen, Ultraschall, CT, Nuklearmedizin
3
In Response
3
Die Virologie von SARS-CoV-2
3
Corona, Immissionen und der Verbrennungsmotor
3
Poster
3
Name: title, dtype: int64

```

The root for such a high amount of title duplicates might be founded for various reasons. At least for this selection, the publications have poor naming which means that they might be published without a name and got assigned a placeholder. Furthermore, this selection has a weak linkage to COVID related research. Based on the titles, the shown selection contains more none COVID19 related publications than research about COVID19.

The column strings contain words that can be investigated by matching certain substrings. Therefore, the following function will support the process.

In [25]:

```

def check_column_for_string(df, col, st):
    """Function which checks matches of substrings in the whole column and returns a dataframe
    """
    df_match = df.loc[df[col].str.contains(st, case=False)]
    return df_match

```

To inspect the density of COVID related publications, the column title will be checked for common words:

- Covid
- Corona
- Lockdown
- SARS

In [26]:

```
matched_df_covid = check_column_for_string(CORD19_CSV, 'title', 'covid')
matched_df_corona = check_column_for_string(CORD19_CSV, 'title', 'corona')
matched_df_lockdown = check_column_for_string(CORD19_CSV, 'title', 'lockdown')
matched_df_sars = check_column_for_string(CORD19_CSV, 'title', 'sars')
```

In [27]:

```
sum_covid_related = len(matched_df_covid) + len(matched_df_corona) + len(matched_df_lockdown)
sum_covid_related
```

Out[27]:

40233

The sum may contain an overlap of the checked substrings which needs to be deducted.

In [28]:

```
matched_df_covid_medical = matched_df_covid.append(matched_df_corona)
matched_df_covid_medical.append(matched_df_lockdown)
matched_df_covid_medical.append(matched_df_sars)
matched_df_covid_medical_len = len(matched_df_covid_medical)
matched_df_covid_medical_len
```

Out[28]:

30982

Furthermore, it is interesting to see the percentage of publications of the CORD19 dataset which match the chosen terms.

In [29]:

```
match_ratio = matched_df_covid_medical_len/len(CORD19_CSV)
match_ratio
```

Out[29]:

0.4000361532899494

To inspect the potential impacts of COVID on society, the column title will be checked for additional terms:

- Economic
- Social

In [30]:

```
matched_df_economic = check_column_for_string(CORD19_CSV, 'title', 'economic')
matched_df_social = check_column_for_string(CORD19_CSV, 'title', 'social')
```

In [31]:

```
len(matched_df_economic)
```

Out[31]:

465

In [32]:

```
len(matched_df_social)
```

Out[32]:

1542

Subsequently, the ratio between the term social and economic is presented.

In [33]:

```
len(matched_df_social)/len(matched_df_economic)
```

Out[33]:

3.3161290322580643

There might be an overlap between the covid related DataFrame and the DataFrame addressing covid impacts on society.

In [34]:

```
matched_df_society = matched_df_economic.append(matched_df_social)
matched_df_society_len = len(matched_df_society)
matched_df_society_len
```

Out[34]:

2007

Compared to the size of the dataset, potential covid impacts on society are more researched and published.

In [35]:

```
combined_df_fields = matched_df_society.append(matched_df_covid_medical)
dif_len_fields = len(matched_df_covid_medical) + len(matched_df_society) - len(combined_df_)
dif_len_fields
```

Out[35]:

0

Both fields do not have any overlap. Potentially, both fields could not be linked to COVID related subjects and investigate distant fields.

Checking the column "source_x"

For publishing purposes, the source of the papers will be inspected.

In [36]:

```
source_x_counted = source_x.value_counts()
source_x_counted
```

Out[36]:

Medline; PMC	32777
Elsevier; Medline; PMC	15500
PMC	15103
MedRxiv; WHO	4760
ArXiv	3115
Elsevier; PMC	1831
BioRxiv; WHO	1252
BioRxiv	796
MedRxiv	367
MedRxiv; Medline; PMC; WHO	331
Medline; PMC; WHO	285
BioRxiv; Medline; PMC; WHO	278
MedRxiv; Medline; WHO	152
BioRxiv; MedRxiv	149
BioRxiv; Medline; WHO	145
Elsevier	135
ArXiv; Medline; PMC	123
Elsevier; Medline; PMC; WHO	115
ArXiv; Elsevier; Medline; PMC	62
BioRxiv; MedRxiv; WHO	51
BioRxiv; Medline; PMC	27
ArXiv; Medline	23
BioRxiv; Medline	13
MedRxiv; Medline; PMC	11
ArXiv; PMC	8
PMC; WHO	7
Elsevier; Medline	6
BioRxiv; MedRxiv; Medline; WHO	6
ArXiv; Elsevier; PMC	5
ArXiv; Elsevier	3
BioRxiv; MedRxiv; Medline; PMC	3
BioRxiv; MedRxiv; Medline; PMC; WHO	3
Elsevier; PMC; WHO	2
MedRxiv; Medline	2
BioRxiv; MedRxiv; Medline	1
ArXiv; Elsevier; Medline; PMC; WHO	1

Name: source_x, dtype: int64

Most of the publications have various and shared publishers.

In [37]:

```
len_source_x_counted = len(source_x_counted)
```

Thusly, it needs to be considered which sources are responsible for most publications.

In [38]:

```
i = 0
while i < len_source_x_counted:
    source_x_counted.index.values[i] = source_x_counted.index.values[i].replace(';', '')
    i = i+1
source_x_counted
```

Out[38]:

Medline PMC	32777
Elsevier Medline PMC	15500
PMC	15103
MedRxiv WHO	4760
ArXiv	3115
Elsevier PMC	1831
BioRxiv WHO	1252
BioRxiv	796
MedRxiv	367
MedRxiv Medline PMC WHO	331
Medline PMC WHO	285
BioRxiv Medline PMC WHO	278
MedRxiv Medline WHO	152
BioRxiv MedRxiv	149
BioRxiv Medline WHO	145
Elsevier	135
ArXiv Medline PMC	123
Elsevier Medline PMC WHO	115
ArXiv Elsevier Medline PMC	62
BioRxiv MedRxiv WHO	51
BioRxiv Medline PMC	27
ArXiv Medline	23
BioRxiv Medline	13
MedRxiv Medline PMC	11
ArXiv PMC	8
PMC WHO	7
Elsevier Medline	6
BioRxiv MedRxiv Medline WHO	6
ArXiv Elsevier PMC	5
ArXiv Elsevier	3
BioRxiv MedRxiv Medline PMC	3
BioRxiv MedRxiv Medline PMC WHO	3
Elsevier PMC WHO	2
MedRxiv Medline	2
BioRxiv MedRxiv Medline	1
ArXiv Elsevier Medline PMC WHO	1

Name: source_x, dtype: int64

In [39]:

```
i = 0
df_source_holder = pd.DataFrame(columns=['source_x', 'counts'])
while i < len_source_x_counted:
    helper = source_x_counted.index[i].split()
    for x in helper:
        data = {'source_x': x, 'counts': source_x_counted[i]}
        df_source_holder = df_source_holder.append(pd.DataFrame([data]))
    i = i + 1

df_source_holder = df_source_holder.reset_index()
df_source_holder = df_source_holder.drop(labels=["index"], axis=1)
df_source_holder
```

Out[39]:

	source_x	counts
0	Medline	32777
1	PMC	32777
2	Elsevier	15500
3	Medline	15500
4	PMC	15500
...
91	ArXiv	1
92	Elsevier	1
93	Medline	1
94	PMC	1
95	WHO	1

96 rows × 2 columns

The DataFrame above contains the counts per each source which must be summed up properly.

In [40]:

```
len_df_source_holder = len(df_source_holder)
sources_series = df_source_holder['source_x'].value_counts()
len_sources_series = len(sources_series.index)

df_total_counts_source_x = pd.DataFrame(columns=['total_counts'], index = sources_series.index)
df_total_counts_source_x ['total_counts'] = 0

i = 0
while i < len_sources_series :
    x = 0
    while x < len_df_source_holder:
        if df_source_holder['source_x'][x] == sources_series.index[i]:
            df_total_counts_source_x ['total_counts'][sources_series.index[i]] = df_total_counts_source_x ['total_counts'] + 1
        x = x + 1
    i = i + 1

df_total_counts_source_x_sorted = df_total_counts_source_x.sort_values(by="total_counts", ascending=False)
```


In [41]:

```
x = df_total_counts_source_x_sorted.index
y = df_total_counts_source_x_sorted['total_counts'].values

plt.bar(x,y)
plt.axhline(y=len(CORD19_CSV), linewidth=2, linestyle='--', color='r')

total_publications, = plt.plot(len(CORD19_CSV), label='Total publications', linestyle='--', color='r')
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)

plt.title("Accumulated count of sources")
plt.xlabel("Name of sources")
plt.ylabel("Count")
plt.show()
```



Overall, PMC and Medline participate mostly as a source within this dataset. It is important to obey that several sources can contribute to one publication.

Checking the column "license"

For the column license, it might be interesting to create an overview of all used licenses.

In [42]:

```
license_counted = license.value_counts()  
license_counted
```

Out[42]:

```
no-cc          29395  
cc-by          23384  
els-covid      7374  
medrxiv        5401  
arxiv          3141  
cc-by-nc       2947  
biorxiv        2325  
cc-by-nc-nd    1860  
green-oa       385  
bronze-oa      352  
cc-by-nc-sa    349  
cc0            312  
cc-by-nd       130  
hybrid-oa      71  
gold-oa        20  
cc-by-sa       2  
Name: license, dtype: int64
```

For the publications, there are several licenses used.

Checking the column "publish_time"

For the column publish_time, it might be interesting to present periods when most of the research was published.

In [43]:

```
publish_time_counted = publish_time.value_counts()  
publish_time_counted
```

Out[43]:

```
2020-12-31      362  
2020-09-30      346  
2021-01-06      329  
2020-09-18      327  
2020-05-26      310  
...  
2010-10-22       1  
2018-11-11       1  
2008-07-13       1  
2010-11-27       1  
2012-12-22       1  
Name: publish_time, Length: 5599, dtype: int64
```

Due to the given dates, it can be obtained that there is research published before COVID-19. For plotting

purposes, the data is transformed into a DataFrame.

In [44]:

```
df_publish_time_counted = pd.DataFrame(columns=['publish_time', 'count'])
df_publish_time_counted['publish_time'] = publish_time_counted.index
df_publish_time_counted['count'] = publish_time_counted.values
df_publish_time_counted.sort_values(by=['publish_time'])
```

Out[44]:

	publish_time	count
4736	1962-09-01	1
4535	1964-01-01	1
4712	1964-02-01	1
5162	1968-07-01	1
4068	1970-01-01	2
...
4621	2021-06-15	1
2445	2021-06-30	4
4011	2021-07-13	2
2977	2021-11-30	3
2066	2021-12-31	5

5599 rows × 2 columns

In [45]:

```
test_def = df_publish_time_counted.sort_values(by=['publish_time'])
test_def
```

Out[45]:

	publish_time	count
4736	1962-09-01	1
4535	1964-01-01	1
4712	1964-02-01	1
5162	1968-07-01	1
4068	1970-01-01	2
...
4621	2021-06-15	1
2445	2021-06-30	4
4011	2021-07-13	2
2977	2021-11-30	3
2066	2021-12-31	5

5599 rows × 2 columns

To distinguish research between COVID-related and non-COVID-related publications, two timestamps are considered. For COVID-related research, publish dates after 2019/12/31 are taken into account. Whereas, publish dates for non-COVID-related research investigates publications before 2020/1/1.

In [46]:

```
date_before_2020 = pd.Timestamp(2020, 1, 1)
date_after_2020 = pd.Timestamp(2019, 12, 31)
df_before_2020 = CORD19_CSV[CORD19_CSV['publish_time'] < date_before_2020]
df_after_2020 = CORD19_CSV[CORD19_CSV['publish_time'] > date_after_2020]
```

To see which periods have more publications, the length of the corresponding DataFrame is shown.

As a result, most of the publications are published after 2019/31/12 which means that the dataset contains mostly research linked to COVID-19. Nevertheless, there is a considerable amount of publications before COVID-19.

In [47]:

```
count_before_2020 = len(df_before_2020)
```

In [48]:


```
count_after_2020 = len(df_after_2020)
```

Therefore, a pie chart is plotted to present the ratio.

In [49]:

```
y = np.array([count_before_2020, count_after_2020])


plt.title("Dominance of COVID19 on the dataset")
pie_labels = ["before 2020/01/01 - " + str(count_before_2020) + " entries \n" + str(round(c
    "after 2019/12/31 - " + str(count_after_2020) + " entries \n" + str(round(cou
plt.pie(y, labels = pie_labels)
plt.show()
```



For this, all publication dates are shown within the next plot.

In [50]:

```
df_publish_time_counted.plot(x='publish_time', y='count', kind='line')
plt.title("Publish period of the CORD19 dataset")
plt.xlabel("Year")
plt.ylabel("Count")
plt.show()
```



Due to this plot, several observations are created.

- The publications are starting from the 1960s which means that there were already other SARS viruses under the investigation of researchers.
- The plot shows four swings with varying extents. Therefore, the four major swings can be affiliated with corresponding SARS-pandemics.
 1. ~1995
 2. ~2001
 3. 2005-2010
 4. 2020
- The extent of the swings seems to have a relationship accordingly to the impact of the occurring pandemic. More widespread and global SARS outbreaks seem to indicate a high count of publications. Whereas, regional and smaller outbreaks are reflected with fewer publications.

In [51]:

```
threshold_source_x = len(CORD19_CSV)
threshold_source_x
```

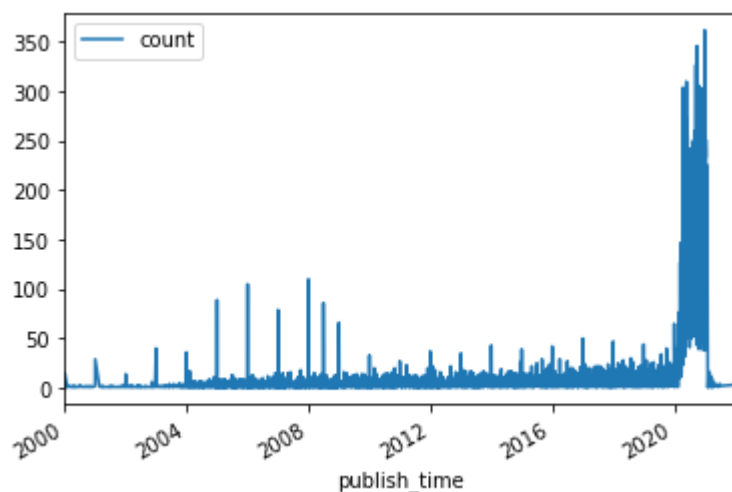
Out[51]:

77448

For further investigation, the next plot will target the last two decades.

In [52]:


```
df_publish_time_counted.plot(x='publish_time', y='count', kind='line', xlim=(pd.Timestamp(2000), pd.Timestamp(2021)), plt.show())
```



In the plot above, two different SARS-outbreak can be seen. In the first decade of this century, the SARS outbreak seems to have less impact because less research was published. The latest SARS outbreak surpassed previous ones by impact and intensity which can be seen in the next plot.

In [53]:

```
df_publish_time_counted.plot(x='publish_time', y='count', kind='line', xlim=(pd.Timestamp(2018), pd.Timestamp(2021)), plt.show())
```



Checking the column "journal"

For this inspection, the NaN values of the column "journal" need to be considered.

In [54]:

```
find_shared_values(journal)
```

Out[54]:

60166

Most of the values of the column are shared ones. It is important to know that each row of the DataFrame is assigned to one journal or none. Thusly, the distribution of journals will be investigated.

In [55]:

```
journal_counted = journal.value_counts()  
journal_counted.head(20)
```

Out[55]:

bioRxiv	2726
PLoS One	2268
Int J Environ Res Public Health	1154
Sci Rep	957
Viruses	613
Virology	424
BMC Infect Dis	421
J Med Virol	394
Arch Virol	386
Emerg Infect Dis	383
PLoS Pathog	356
BMJ Open	352
medRxiv	344
Virol J	333
Front Psychol	326
Clin Infect Dis	316
Computational Science - ICCS 2020	296
Int J Mol Sci	289
Front Immunol	286
BMC Public Health	271

Name: journal, dtype: int64

Compared to other columns, the distribution of journals is equally.

Checking the column "url"

The column "url" will be investigated to present contributing organisations by hostnames.

In [56]:

```
url_counted = url.value_counts()
url_counted
```

Out[56]:

```
https://api.elsevier.com/content/article/pii/S1386653220303589; (https://api.elsevier.com/content/article/pii/S1386653220303589;) https://www.ncbi.nlm.nih.gov/pubmed/32891938/; (https://www.ncbi.nlm.nih.gov/pubmed/32891938/;) https://www.sciencedirect.com/science/article/pii/S1386653220303589?v=s5; (https://www.sciencedirect.com/science/article/pii/S1386653220303589?v=s5;) https://doi.org/10.1016/j.jcv.2020.104616 (https://doi.org/10.1016/j.jcv.2020.104616)      1
https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7686826/ (https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7686826/)
1
https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7286207/ (https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7286207/)
1
http://medrxiv.org/cgi/content/short/2020.07.27.20159996v1?rss=1; (http://medrxiv.org/cgi/content/short/2020.07.27.20159996v1?rss=1;) https://doi.org/10.1101/2020.07.27.20159996 (https://doi.org/10.1101/2020.07.27.20159996)
1
https://www.ncbi.nlm.nih.gov/pubmed/32324273/; (https://www.ncbi.nlm.nih.gov/pubmed/32324273/;) https://doi.org/10.1002/cncr.32919 (https://doi.org/10.1002/cncr.32919)
1
..
https://doi.org/10.1002/jdd.12364; (https://doi.org/10.1002/jdd.12364;) https://www.ncbi.nlm.nih.gov/pubmed/32813876/ (https://www.ncbi.nlm.nih.gov/pubmed/32813876/)
1
https://arxiv.org/pdf/2005.13754v1.pdf (https://arxiv.org/pdf/2005.13754v1.pdf)
1
https://doi.org/10.1371/journal.pone.0242474; (https://doi.org/10.1371/journal.pone.0242474;) https://www.ncbi.nlm.nih.gov/pubmed/33216795/ (https://www.ncbi.nlm.nih.gov/pubmed/33216795/)
1
https://doi.org/10.1186/s13063-020-04997-6; (https://doi.org/10.1186/s13063-020-04997-6;) https://www.ncbi.nlm.nih.gov/pubmed/33509275/ (https://www.ncbi.nlm.nih.gov/pubmed/33509275/)
1
https://www.ncbi.nlm.nih.gov/pubmed/32238612/ (https://www.ncbi.nlm.nih.gov/pubmed/32238612/)
1
Name: url, Length: 77448, dtype: int64
```

Some entries contain more than one URL pointing to a publication. Consequently, entries with more than one URL are split into a series containing one URL per tuple.

In [57]:

```
%%time
CORD19_CSV.url = CORD19_CSV.url.str.replace(';', '')
url = CORD19_CSV.url
url_list= []
for row in url:
    row_string = re.split('\s+', row)
    for i in row_string:
        url_list.append(urlparse(i).netloc)
url_list
```

Wall time: 9.09 s

Next, all URLs are counted to see which host publishes research and to which extent. For this, it is crucial to know that the same publication is accessible via different hosts.

In [58]:

```
url_series = pd.Series(url_list)
url_counted = url_series.value_counts()
url_counted.head(20)
```

Out[58]:

www.ncbi.nlm.nih.gov	64918
doi.org	54410
api.elsevier.com	17660
www.sciencedirect.com	17660
medrxiv.org	4995
arxiv.org	3340
europemc.org	313
academic.oup.com	118
www.cell.com	46
link.springer.com	42
onlinelibrary.wiley.com	36
www.nature.com	23
jvi.asm.org	16
journals.iucr.org	14
www.jbc.org	14
www.thelancet.com	8
pubs.acs.org	7
www.tandfonline.com	6
www.biorxiv.org	4
ajp.amjpathol.org	3

dtype: int64

As a result, the National Center for Biotechnology Information (NCBI) is responsible for almost all publications. This means that the publications are hosted by different organisations holding and securing the research as a possible backup.

Checking the column "software"

The analysis of the column software is conducted in the Jupyter Notebook named "CORD-19-software-counting-cs5099".